

Cuisine and Yelp Ratings*

Moohaeng Sohn

April 18, 2024

First sentence. Second sentence. Third sentence. Fourth sentence.

Table of contents

1	Introduction	2
2	Data	2
2.1	Star Rating	3
2.2	Categories (Cuisine)	4
2.3	Price Range	4
2.4	Measurement	5
3	Model	5
3.1	Model set-up	6
3.1.1	Model justification	6
4	Results	7
5	Discussion	7
5.1	First discussion point	7
5.2	Second discussion point	8
5.3	Third discussion point	8
5.4	Weaknesses and next steps	8
	Appendix	9
A	Model details	9
A.1	Posterior predictive check	9
A.2	Diagnostics	9

*Code and data are available at: <https://github.com/alexsohn1126/yelp-analysis>

1 Introduction

Thanks to internet technology, we can get hundreds, if not thousands of reviews of restaurants around us. This means we can choose which restaurants we will go to based on those reviews. Therefore, keeping a high review rating is very important for a restaurant's long-term success.

Yelp is a website where people can post reviews about local businesses, containing 287 million reviews (n.d.a). Yelp's reviews are based on the 5-star system. Users can choose between one to five stars to put on their review. These reviews are collected and averaged, which becomes the rating for the establishment. We will use a dataset from Yelp which we will dive into in Section 2.

TODO: ADD THINGS ABOUT WHAT WE FOUND OUT AND WHY IS IT IMPORTANT

In this paper, our estimand of interest is how different factors such as the cuisine of the restaurant, or the price of the menu items in the restaurants affect the rating of a restaurant. We will first explore the dataset in Section 2, and discuss the model the relationship between aforementioned variables using logistic regression in Section 3. Then we will look at the results in Section 4, finally discussing about these results in Section 5.

We used the programming language R (R Core Team 2023), along with packages `tidyverse` (Wickham et al. 2019), `rstanarm` (Goodrich et al. 2022), `jsonlite` (Ooms 2014), `arrow` (Richardson et al. 2024), `modelsummary` (Arel-Bundock 2022), `kableExtra` (Zhu 2021), `here` (Müller 2020).

2 Data

The restaurant review dataset we will use is from Yelp. Yelp offers an academic dataset for the public to use, although it is a small subset of their massive database (n.d.b). The dataset is split between multiple JSON files, but we will only focus on businesses JSON. The raw businesses dataset contains 150,346 businesses. These businesses are from metropolitan areas in United States and Canada. Specifically, from metropolitan areas near Montreal, Calgary, Toronto, Pittsburgh, Charlotte, Urbana-Champaign, Phoenix, Las Vegas, Madison, and Cleveland (n.d.b).

Another possible source of restaurant reviews could have been from Google reviews of restaurants, but that requires us to use Google Business Profile APIs, which cost money. There is always an option to perform webscraping, but this is a gray area legally, and may be computationally expensive. There are other review websites such as TripAdvisor, but they do not seem to have a dataset open to the public like Yelp does.

The dataset contains multiple variables such as the name of the business, the location of the business, and things such as amenities on site. We have filtered through the raw dataset to only contain restaurants. There were a total of 47,213 restaurants in the final dataset. We will focus on the star rating, categories, and price range of the menu in the restaurant.

2.1 Star Rating

Star rating is the average star rating of a restaurant. One odd thing about the star rating given in the dataset is that it is rounded to the nearest star or half of a star. So all possible values of star ratings are: 1.0, 1.5, 2.0, and so on until 5.0.

Figure 1 shows us that 4.0 is the most common star rating of restaurants. This distribution is left skewed, as we can see the rating gradually increases from 1.0 to 4.0, then quickly drops off from 4.0 to 5.0. We can infer from this that most people consider somewhere around 4-stars an average dining experience.

Table 1 shows the summary statistics of star ratings, and we can see that as we explained, that the minimum is 1.0 and the maximum is 5.0. The standard deviation isn't that big, which is to be expected because most of the restaurants lie between 3.5 to 4.5 rating range.




Figure 1: Distribution of Star Ratings

Table 1: Summary Statistics of Star Ratings

Count	Mean	Median	SD	Minimum	Maximum
47213	3.482187	3.5	0.8082359	1	5

Table 2: Numbers of Restaurants Per Cuisine

Cuisine	Number of Restaurants	Percentage (%)
Other	21741	46.05
American	10841	22.96
Italian	3729	7.90
Mexican	3476	7.36
Chinese	2423	5.13
Japanese	1351	2.86
Mediterranean	738	1.56
Thai	694	1.47
Indian	628	1.33
Vietnamese	617	1.31
Caribbean	343	0.73
French	323	0.68
Middle Eastern	309	0.65

2.2 Categories (Cuisine)

Categories for a restaurant describes what kind of business it is. These categories are chosen manually by the business owners (2024). Because Yelp is a business review platform, we have decided to filter out non-restaurant businesses from our final dataset. This meant any restaurants which didn’t include one of the categories: “Restaurants”, “Food”, “Fast Food” would not be included in the final dataset. After filtering out non-restaurants, we have chosen top 12 cuisines to categorize each restaurant into. The restaurants that did not have these cuisines in their categories were put into “Other” category. This is why we have named this section also cuisine, as we will focus on cuisine categories of restaurants.

Table 2 shows us what types of cuisines we have selected, and the number of restaurants with that cuisine in their category. We can see nearly half of the restaurants are categorized as “Other”. The next largest cuisine in our dataset is American. This may be due to the fact that all the restaurants in the dataset are from United States and Canada.

2.3 Price Range

The price range is given as an integer from 1 to 4, inclusive. This is supposed to indicate how expensive the restaurant is per person. 1 being more affordable, and 4 being more expensive.

Table 3: Numbers of Restaurants Per Price Range

Price Range	Number of Restaurants	Percentage (%)
1	20670	43.78
2	24566	52.03
3	1762	3.73
4	215	0.46

There were restaurants with no price ranges, and those were discarded from the cleaned dataset. On Yelp’s website, these price ranges are shown as number of dollar signs equal to the price range. For example, Yelp would show one dollar sign if price range is 1.

Table 3 shows that the over 95% of the restaurants have price range of 1 or 2. About 95.81% of restaurants in the database are in price range 1 or 2. Only 3.73% of the restaurants are in price range 3, and even less, 0.46% of the restaurants are in price range 4.

2.4 Measurement

For star ratings, the minimum star rating is one and the maximum star rating is five. This is because the minimum star rating that someone can give a restaurant is 1 stars, and the maximum star rating is 5 stars. While the average star rating can have 0.5 star increments, users’ reviews can only give ratings in 1 star increments. For example, no 4.5 stars can be given in a review. Each user can rate a business only once, though they are free to change their rating later, one user cannot post multiple reviews about a restaurant. This limits the power each user has to change the rating of a business.

As mentioned above, categories for a restaurant is not given automatically, rather given by the business owners. This means it can be a bit inaccurate when it comes to what exact cuisine that a restaurant serves. Another way the cuisines can be inaccurate is when a restaurant serves foods from multiple cuisines. This makes it complicated to categorize them exactly into each cuisine.

The exact method of how the price range was decided by Yelp is not known. We can make guesses about what kind of information they use to determine the price range, but no clear documentation is available for the price range statistics. We were able to find online discussions about the price range being decided by the check-in feature of Yelp, possibly implying that there are certain price ranges that Yelp uses to display these price ranges (S 2017).

3 Model

We are trying to model the rating y , and as we have discussed in the data section, there are only 9 possible outcomes. From 1.0, to 5.0, in increments of 0.5. Therefore we will use ordinal

logistic regression. If we assume that the observed rating y is from a continuous variable y^* , then we can say that:

$$y = \begin{cases} 1.0 & \text{if } y^* \leq c_{1.0|1.5} \\ 1.5 & \text{if } y^* \in (c_{1.0|1.5}, c_{1.5|2.0}) \\ \vdots & \\ 4.5 & \text{if } y^* \in (c_{4.0|4.5}, c_{4.5|5.0}) \\ 5.0 & \text{if } y^* \geq c_{4.5|5.0} \end{cases}$$

Where $c_{n_1|n_2}$ represents the cutoff between n_1 rating and n_2 rating. The above notations are borrowed from Gelman, Hill, and Vehtari (2020) [p.276] and Alexander et al. (2024, 12).

We will use a Bayesian framework to make our model. To do this, we use **rstanarm** (Goodrich et al. 2022) R package. Background details and diagnostics are included in Appendix A.

3.1 Model set-up

We will model the final discrete rating y by a continuous variable y^* . Because we are doing ordinal logistic regression, y^* will have a logistic distribution.

$$y^* = \beta_1 \cdot \text{cuisine}_i + \beta_2 \cdot \text{price}_i$$

We use the R2 prior, because **rstanarm** requires us to do this for this computational process, with the mean set as 0.3 as suggested by Gelman, Hill, and Vehtari (2020, 276). R2 uses the prior beliefs about the location about R^2 , which is “the proportion of variance in the outcome attributable to the predictors” Goodrich et al. (2022). We use default priors from **rstanarm** for the cut points ($c_{n_1|n_2}$), which is a dirichlet distribution with concentration of 1.

3.1.1 Model justification

We can expect that if a restaurant has higher prices, then the restaurant will be higher quality. This should be reflected in the ratings by the customers, and those restaurants with higher quality will have higher ratings. However, one could set much higher expectations for more expensive restaurants, effectively cancelling out the higher quality food and service that is expected from a higher priced restaurant. We will keep this in mind when we discuss our results.

We also expect that different cuisines will have different ratings. This dataset is exclusively about restaurants in North America. Therefore we may expect that reviewers are more likely to be from North America, so perhaps they would give a higher rating to a restaurant that serves western cuisine, such as Italian, or American.

4 Results

Our results are summarized in `?@tbl-modelresults`, and `?@fig-modelresults`.

The rows with numbers and a vertical line between them are the cutpoints. For example, a rating would be considered 1.5 if we had a y^* that is greater than 1.0|1.5 and less than 1.5|2.0.

We can see that nearly all cuisines has a positive effect on the rating, except for Chinese.

5 Discussion

5.1 First discussion point

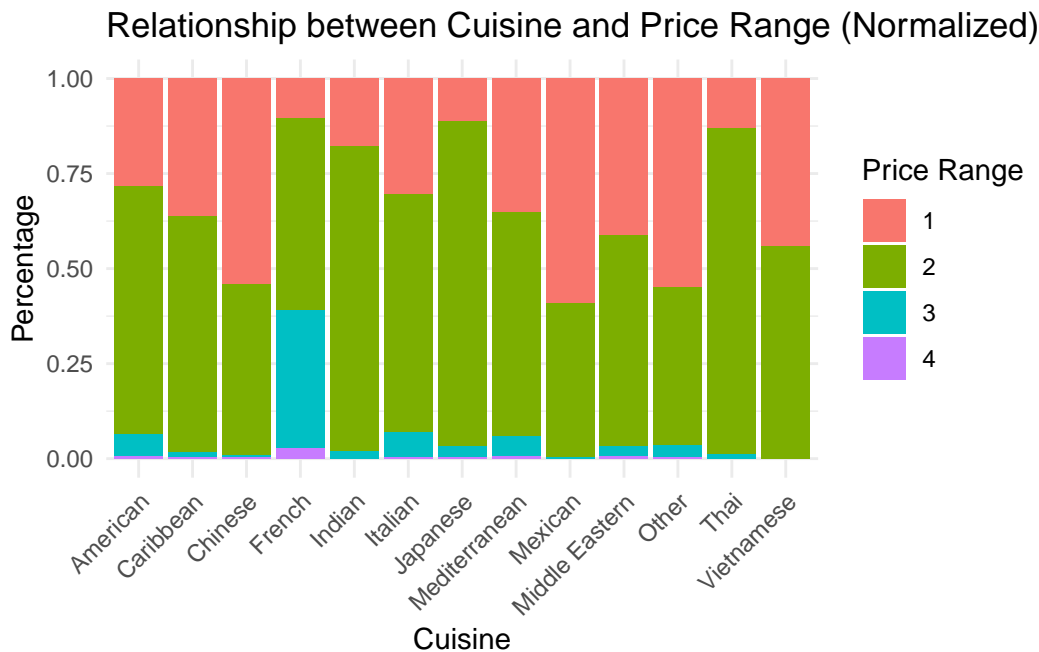


Figure 2: Bills of penguins

If my paper were 10 pages, then should be be at least 2.5 pages. The discussion is a chance to show off what you know and what you learnt from all this.

5.2 Second discussion point

5.3 Third discussion point

5.4 Weaknesses and next steps

Weaknesses and next steps should also be included.

Appendix

A Model details

A.1 Posterior predictive check

In Figure 3 we implement a posterior predictive check. This shows that there is a large disparity between the predictions and the actual data.

In Figure 4 we compare the posterior with the prior. This shows the posteriors are around where the credible intervals are in priors.

A.2 Diagnostics

Figure 5 is a trace plot. It shows the chains cover a constant amount of region. This suggests that a large amount of sample is chosen at around the target space we want to explore.

Figure 6 is a Rhat plot. It shows all of the Rhat values are very close to 1, meaning the chains had enough time to converge.




Figure 3: Examining how the model fits, and is affected by the data - Posterior prediction check




Figure 4: Examining how the model fits, and is affected by the data - Comparing the posterior with the prior




Figure 5: Checking the convergence of the MCMC algorithm - Trace Plot



Figure 6: Checking the convergence of the MCMC algorithm - Rhat Plot

References

2024. *Yelp*. https://blog.yelp.com/businesses/yelp_category_list/.
- . n.d.a. *Yelp - Company - Fast Facts*. <https://www.yelp-press.com/company/fast-facts/default.aspx>.
- . n.d.b. *Yelp Dataset FAQ*. <https://www.yelp.com/dataset/documentation/faq>.
- Alexander, Rohan, Lindsay Katz, Callandra Moor, Michaela Drouillard, Michael Wing-Cheung Wong, and Zane Schwartz. 2024. “Evaluating the Decency and Consistency of Data Validation Tests Generated by LLMs.” https://github.com/RohanAlexander/evaluating_decency_and_consistency.
- Arel-Bundock, Vincent. 2022. “modelssummary: Data and Model Summaries in R.” *Journal of Statistical Software* 103 (1): 1–23. <https://doi.org/10.18637/jss.v103.i01>.
- Gelman, Andrew, Jennifer Hill, and Aki Vehtari. 2020. *Regression and Other Stories*. Cambridge University Press. <https://avehtari.github.io/ROS-Examples/>.
- Goodrich, Ben, Jonah Gabry, Imad Ali, and Sam Brilleman. 2022. “Rstanarm: Bayesian Applied Regression Modeling via Stan.” <https://mc-stan.org/rstanarm/>.
- Müller, Kirill. 2020. *Here: A Simpler Way to Find Your Files*. <https://CRAN.R-project.org/package=here>.
- Ooms, Jeroen. 2014. “The Jsonlite Package: A Practical and Consistent Mapping Between JSON Data and r Objects.” *arXiv:1403.2805 [Stat.CO]*. <https://arxiv.org/abs/1403.2805>.
- R Core Team. 2023. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Richardson, Neal, Ian Cook, Nic Crane, Dewey Dunnington, Romain François, Jonathan Keane, Dragoş Moldovan-Grünfeld, Jeroen Ooms, Jacob Wujciak-Jens, and Apache Arrow. 2024. *Arrow: Integration to 'Apache' 'Arrow'*. <https://CRAN.R-project.org/package=arrow>.
- S, Dez. 2017. *Yelp*. <https://www.yelp.com/topic/mesa-what-is-latest-definition-and-range-of-the-dollar-sign-on-yelp-for-actual-expense>.
- Wickham, Hadley, Mara Averick, Jennifer Bryan, Winston Chang, Lucy D’Agostino McGowan, Romain François, Garrett Golemund, et al. 2019. “Welcome to the tidyverse.” *Journal of Open Source Software* 4 (43): 1686. <https://doi.org/10.21105/joss.01686>.
- Zhu, Hao. 2021. *kableExtra: Construct Complex Table with 'Kable' and Pipe Syntax*. <https://CRAN.R-project.org/package=kableExtra>.